

# 601 Individual Assignment

Olalekan Fagbuyi

2023-10-10

```
setwd('C:\\Users\\User\\Desktop\\WLU')
```

## 1. Business Understanding

This project aims to predict the popularity of online articles using the number of shares as a metric for measuring popularity. Correctly classifying these online articles is important to the company because it implies we will be able to optimize revenue regeneration (\$0.75 for the first 1000 shares and \$2 after) considering our 12 articles per day limit constraint.

A threshold of 1400 has been set to determine which articles are popular and which are not. Shares Value Range: Number of Instances in Range: < 1400 18490 (46.64%) and >= 1400 21154 (53.35%).

## 2. Data Understanding

The dataset contain 39644 articles with 61 attributes will be analyzed and then used to build a machine model for making predictions.

Dataset Breakdown: 61 attributes (58 predictive attributes, 2 non-predictive, 1 goal field)

### 2.1 Importing Libraries and Loading Dataset

```
#importing library
library(caret)

## Loading required package: ggplot2

## Loading required package: lattice

library(tidyverse)

## — Attaching core tidyverse packages ————— tidyverse
2.0.0 —
## ✓ dplyr      1.1.3      ✓ readr      2.1.4
## ✓ forcats   1.0.0      ✓ stringr   1.5.0
## ✓ lubridate 1.9.2      ✓ tibble    3.2.1
## ✓ purrr     1.0.2      ✓ tidyr     1.3.0

## — Conflicts —————
tidyverse_conflicts() —
## ✗ dplyr::filter() masks stats::filter()
## ✗ dplyr::lag()    masks stats::lag()
## ✗ purrr::lift()   masks caret::lift()
```

```
## i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all  
conflicts to become errors
```

```
library(randomForest)
```

```
## randomForest 4.7-1.1  
## Type rfNews() to see new features/changes/bug fixes.  
##  
## Attaching package: 'randomForest'  
##  
## The following object is masked from 'package:dplyr':  
##  
##     combine  
##  
## The following object is masked from 'package:ggplot2':  
##  
##     margin
```

```
library(corrplot)
```

```
## corrplot 0.92 loaded
```

```
library(rpart)
```

```
library(kernlab)
```

```
##  
## Attaching package: 'kernlab'  
##  
## The following object is masked from 'package:purrr':  
##  
##     cross  
##  
## The following object is masked from 'package:ggplot2':  
##  
##     alpha
```

```
library(pROC)
```

```
## Type 'citation("pROC")' for a citation.  
##  
## Attaching package: 'pROC'  
##  
## The following objects are masked from 'package:stats':  
##  
##     cov, smooth, var
```

```
library(gridExtra)
```

```
##  
## Attaching package: 'gridExtra'  
##  
## The following object is masked from 'package:randomForest':
```

```
##
##      combine
##
## The following object is masked from 'package:dplyr':
##
##      combine

#Loading dataset
articles <- read.csv("OnlineNewsPopularity.csv")
head(articles)

##                                     url timedelta
## 1  http://mashable.com/2013/01/07/amazon-instant-video-browser/      731
## 2  http://mashable.com/2013/01/07/ap-samsung-sponsored-tweets/      731
## 3  http://mashable.com/2013/01/07/apple-40-billion-app-downloads/      731
## 4  http://mashable.com/2013/01/07/astronaut-notre-dame-bcs/          731
## 5  http://mashable.com/2013/01/07/att-u-verse-apps/                  731
## 6  http://mashable.com/2013/01/07/beewi-smart-toys/                  731
##  n_tokens_title n_tokens_content n_unique_tokens n_non_stop_words
## 1             12             219         0.6635945             1
## 2              9             255         0.6047431             1
## 3              9             211         0.5751295             1
## 4              9             531         0.5037879             1
## 5             13            1072         0.4156456             1
## 6             10             370         0.5598886             1
##  n_non_stop_unique_tokens num_hrefs num_self_hrefs num_imgs num_videos
## 1             0.8153846           4           2           1           0
## 2             0.7919463           3           1           1           0
## 3             0.6638655           3           1           1           0
## 4             0.6656347           9           0           1           0
## 5             0.5408895          19          19          20           0
## 6             0.6981982           2           2           0           0
##  average_token_length num_keywords data_channel_is_lifestyle
## 1             4.680365           5           0
## 2             4.913725           4           0
## 3             4.393365           6           0
## 4             4.404896           7           0
## 5             4.682836           7           0
## 6             4.359459           9           0
##  data_channel_is_entertainment data_channel_is_bus data_channel_is_socmed
## 1              1              0              0
## 2              0              1              0
## 3              0              1              0
## 4              1              0              0
## 5              0              0              0
## 6              0              0              0
##  data_channel_is_tech data_channel_is_world kw_min_min kw_max_min
## 1              0              0              0              0
0
```

```

## 2          0          0          0          0
0
## 3          0          0          0          0
0
## 4          0          0          0          0
0
## 5          1          0          0          0
0
## 6          1          0          0          0
0
##   kw_min_max kw_max_max kw_avg_max kw_min_avg kw_max_avg kw_avg_avg
## 1          0          0          0          0          0          0
## 2          0          0          0          0          0          0
## 3          0          0          0          0          0          0
## 4          0          0          0          0          0          0
## 5          0          0          0          0          0          0
## 6          0          0          0          0          0          0
##   self_reference_min_shares self_reference_max_shares
## 1                   496                   496
## 2                   0                   0
## 3                   918                   918
## 4                   0                   0
## 5                   545                  16000
## 6                  8500                  8500
##   self_reference_avg_sharess weekday_is_monday weekday_is_tuesday
## 1                   496.000                   1                   0
## 2                   0.000                   1                   0
## 3                   918.000                   1                   0
## 4                   0.000                   1                   0
## 5                  3151.158                   1                   0
## 6                  8500.000                   1                   0
##   weekday_is_wednesday weekday_is_thursday weekday_is_friday
## 1                   0                   0                   0
## 2                   0                   0                   0
## 3                   0                   0                   0
## 4                   0                   0                   0
## 5                   0                   0                   0
## 6                   0                   0                   0
##   weekday_is_saturday weekday_is_sunday is_weekend   LDA_00   LDA_01
## 1                   0                   0         0 0.50033120 0.37827893
## 2                   0                   0         0 0.79975569 0.05004668
## 3                   0                   0         0 0.21779229 0.03333446
## 4                   0                   0         0 0.02857322 0.41929964
## 5                   0                   0         0 0.02863281 0.02879355
## 6                   0                   0         0 0.02224528 0.30671758
##   LDA_02   LDA_03   LDA_04 global_subjectivity
## 1 0.04000468 0.04126265 0.04012254         0.5216171
## 2 0.05009625 0.05010067 0.05000071         0.3412458
## 3 0.03335142 0.03333354 0.68218829         0.7022222
## 4 0.49465083 0.02890472 0.02857160         0.4298497

```

```

## 5 0.02857518 0.02857168 0.88542678          0.5135021
## 6 0.02223128 0.02222429 0.62658158          0.4374086
##   global_sentiment_polarity global_rate_positive_words
## 1          0.09256198          0.04566210
## 2          0.14894781          0.04313725
## 3          0.32333333          0.05687204
## 4          0.10070467          0.04143126
## 5          0.28100348          0.07462687
## 6          0.07118419          0.02972973
##   global_rate_negative_words rate_positive_words rate_negative_words
## 1          0.013698630          0.7692308          0.2307692
## 2          0.015686275          0.7333333          0.2666667
## 3          0.009478673          0.8571429          0.1428571
## 4          0.020715631          0.6666667          0.3333333
## 5          0.012126866          0.8602151          0.1397849
## 6          0.027027027          0.5238095          0.4761905
##   avg_positive_polarity min_positive_polarity max_positive_polarity
## 1          0.3786364          0.1000000          0.7
## 2          0.2869146          0.0333333          0.7
## 3          0.4958333          0.1000000          1.0
## 4          0.3859652          0.1363636          0.8
## 5          0.4111274          0.0333333          1.0
## 6          0.3506100          0.1363636          0.6
##   avg_negative_polarity min_negative_polarity max_negative_polarity
## 1          -0.3500000          -0.600          -0.2000000
## 2          -0.1187500          -0.125          -0.1000000
## 3          -0.4666667          -0.800          -0.1333333
## 4          -0.3696970          -0.600          -0.1666667
## 5          -0.2201923          -0.500          -0.0500000
## 6          -0.1950000          -0.400          -0.1000000
##   title_subjectivity title_sentiment_polarity abs_title_subjectivity
## 1          0.5000000          -0.1875000          0.0000000
## 2          0.0000000          0.0000000          0.5000000
## 3          0.0000000          0.0000000          0.5000000
## 4          0.0000000          0.0000000          0.5000000
## 5          0.4545455          0.1363636          0.04545455
## 6          0.6428571          0.2142857          0.14285714
##   abs_title_sentiment_polarity shares
## 1          0.1875000          593
## 2          0.0000000          711
## 3          0.0000000          1500
## 4          0.0000000          1200
## 5          0.1363636          505
## 6          0.2142857          855

## 'data.frame':   39644 obs. of  61 variables:
## $ url          : chr
"http://mashable.com/2013/01/07/amazon-instant-video-browser/"
"http://mashable.com/2013/01/07/ap-samsung-sponsored-tweets/"
"http://mashable.com/2013/01/07/apple-40-billion-app-downloads/"

```

```

"http://mashable.com/2013/01/07/astronaut-notre-dame-bcs/" ...
## $ timedelta : num 731 731 731 731 731 731 731 731 731 731
731 ...
## $ n_tokens_title : num 12 9 9 9 13 10 8 12 11 10 ...
## $ n_tokens_content : num 219 255 211 531 1072 ...
## $ n_unique_tokens : num 0.664 0.605 0.575 0.504 0.416 ...
## $ n_non_stop_words : num 1 1 1 1 1 ...
## $ n_non_stop_unique_tokens : num 0.815 0.792 0.664 0.666 0.541 ...
## $ num_hrefs : num 4 3 3 9 19 2 21 20 2 4 ...
## $ num_self_hrefs : num 2 1 1 0 19 2 20 20 0 1 ...
## $ num_imgs : num 1 1 1 1 20 0 20 20 0 1 ...
## $ num_videos : num 0 0 0 0 0 0 0 0 0 1 ...
## $ average_token_length : num 4.68 4.91 4.39 4.4 4.68 ...
## $ num_keywords : num 5 4 6 7 7 9 10 9 7 5 ...
## $ data_channel_is_lifestyle : num 0 0 0 0 0 0 1 0 0 0 ...
## $ data_channel_is_entertainment : num 1 0 0 1 0 0 0 0 0 0 ...
## $ data_channel_is_bus : num 0 1 1 0 0 0 0 0 0 0 ...
## $ data_channel_is_socmed : num 0 0 0 0 0 0 0 0 0 0 ...
## $ data_channel_is_tech : num 0 0 0 0 1 1 0 1 1 0 ...
## $ data_channel_is_world : num 0 0 0 0 0 0 0 0 0 1 ...
## $ kw_min_min : num 0 0 0 0 0 0 0 0 0 0 ...
## $ kw_max_min : num 0 0 0 0 0 0 0 0 0 0 ...
## $ kw_avg_min : num 0 0 0 0 0 0 0 0 0 0 ...
## $ kw_min_max : num 0 0 0 0 0 0 0 0 0 0 ...
## $ kw_max_max : num 0 0 0 0 0 0 0 0 0 0 ...
## $ kw_avg_max : num 0 0 0 0 0 0 0 0 0 0 ...
## $ kw_min_avg : num 0 0 0 0 0 0 0 0 0 0 ...
## $ kw_max_avg : num 0 0 0 0 0 0 0 0 0 0 ...
## $ kw_avg_avg : num 0 0 0 0 0 0 0 0 0 0 ...
## $ self_reference_min_shares : num 496 0 918 0 545 8500 545 545 0 0
...
## $ self_reference_max_shares : num 496 0 918 0 16000 8500 16000 16000
0 0 ...
## $ self_reference_avg_shareess : num 496 0 918 0 3151 ...
## $ weekday_is_monday : num 1 1 1 1 1 1 1 1 1 1 ...
## $ weekday_is_tuesday : num 0 0 0 0 0 0 0 0 0 0 ...
## $ weekday_is_wednesday : num 0 0 0 0 0 0 0 0 0 0 ...
## $ weekday_is_thursday : num 0 0 0 0 0 0 0 0 0 0 ...
## $ weekday_is_friday : num 0 0 0 0 0 0 0 0 0 0 ...
## $ weekday_is_saturday : num 0 0 0 0 0 0 0 0 0 0 ...
## $ weekday_is_sunday : num 0 0 0 0 0 0 0 0 0 0 ...
## $ is_weekend : num 0 0 0 0 0 0 0 0 0 0 ...
## $ LDA_00 : num 0.5003 0.7998 0.2178 0.0286 0.0286
...
## $ LDA_01 : num 0.3783 0.05 0.0333 0.4193 0.0288
...
## $ LDA_02 : num 0.04 0.0501 0.0334 0.4947 0.0286
...
## $ LDA_03 : num 0.0413 0.0501 0.0333 0.0289 0.0286
...

```

```
## $ LDA_04 : num 0.0401 0.05 0.6822 0.0286 0.8854
...
## $ global_subjectivity : num 0.522 0.341 0.702 0.43 0.514 ...
## $ global_sentiment_polarity : num 0.0926 0.1489 0.3233 0.1007 0.281
...
## $ global_rate_positive_words : num 0.0457 0.0431 0.0569 0.0414 0.0746
...
## $ global_rate_negative_words : num 0.0137 0.01569 0.00948 0.02072
0.01213 ...
## $ rate_positive_words : num 0.769 0.733 0.857 0.667 0.86 ...
## $ rate_negative_words : num 0.231 0.267 0.143 0.333 0.14 ...
## $ avg_positive_polarity : num 0.379 0.287 0.496 0.386 0.411 ...
## $ min_positive_polarity : num 0.1 0.0333 0.1 0.1364 0.0333 ...
## $ max_positive_polarity : num 0.7 0.7 1 0.8 1 0.6 1 1 0.8 0.5 ...
## $ avg_negative_polarity : num -0.35 -0.119 -0.467 -0.37 -0.22 ...
## $ min_negative_polarity : num -0.6 -0.125 -0.8 -0.6 -0.5 -0.4 -
0.5 -0.5 -0.125 -0.5 ...
## $ max_negative_polarity : num -0.2 -0.1 -0.133 -0.167 -0.05 ...
## $ title_subjectivity : num 0.5 0 0 0 0.455 ...
## $ title_sentiment_polarity : num -0.188 0 0 0 0.136 ...
## $ abs_title_subjectivity : num 0 0.5 0.5 0.5 0.0455 ...
## $ abs_title_sentiment_polarity : num 0.188 0 0 0 0.136 ...
## $ shares : int 593 711 1500 1200 505 855 556 891
3600 710 ...
```


## 2.2 Determining Popularity Threshold

In order to predict popularity of articles, an average threshold will be determined from the shares column.

The mean would be used if the column is normally distributed with no outliers. However, if there are outliers, the median of the column will be picked as this measure is less sensitive to extreme values.

```
# Create a box plot for a specific column
boxplot(articles$shares, main = "Box Plot of Shares",
        ylab = "shares")
```

## Box Plot of Shares



```
#No of outliers in the shares column  
outliers <- boxplot.stats(articles$shares)$out  
length(outliers)  
## [1] 4541
```



The shares column has a considerable amount of outliers (4541 or 11.45% of total data points). The median will be used as an average instead of the mean on this occasion.

```
median(articles$shares)
```

```
## [1] 1400
```

### 2.3 Creating popularity column based on Threshold

```
articles$popular <- ifelse(articles$shares >= 1400, 1, 0)
```

*#Dropping shares column*

```
articles1 <- subset( articles, select = -c(shares) )
```

## 3 Data Visualization

Getting a better understanding of the dataset via visualizations

### 3.1 Popularity Distribution

```
ggplot(articles1, aes(x = popular)) +
```

```
geom_bar() +
```


```
geom_text(stat='count', aes(label=..count..), vjust=-1)
```

```
## Warning: The dot-dot notation (`..count..`) was deprecated in ggplot2  
3.4.0.
```

```
## i Please use `after_stat(count)` instead.
```

```
## This warning is displayed once every 8 hours.
```

```
## Call `lifecycle::last_lifecycle_warnings()` to see where this warning was  
## generated.
```

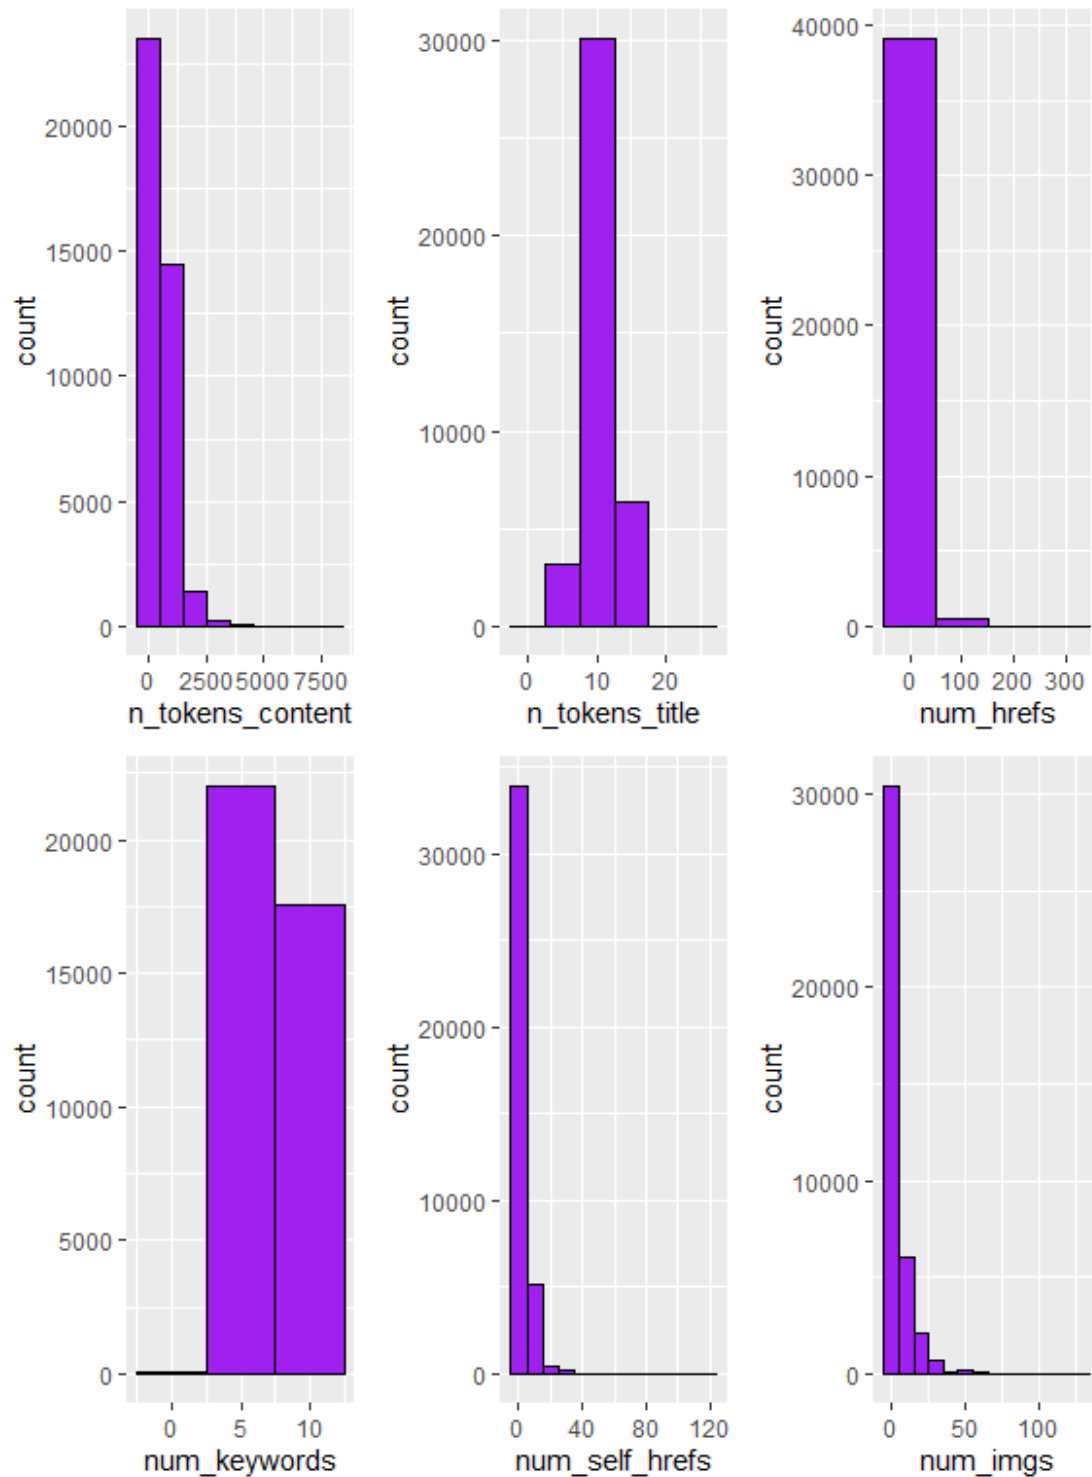


### 3.2 - Univariate Analysis - num columns

```
p1 <- ggplot(articles1) + geom_histogram(aes(n_tokens_content), binwidth =  
1000, fill = "purple", col = "black")  
p2 <- ggplot(articles1) + geom_histogram(aes(n_tokens_title), binwidth = 5,  
fill = "purple", col = "black")
```

```
p3 <- ggplot(articles1) + geom_histogram(aes(num_hrefs), binwidth = 100, fill
="purple",col ="black")
p4 <- ggplot(articles1) + geom_histogram(aes(num_keywords ), binwidth =5,
fill ="purple",col ="black")
p5 <- ggplot(articles1) + geom_histogram(aes(num_self_hrefs), binwidth = 10,
fill ="purple",col ="black")
p6 <- ggplot(articles1) + geom_histogram (aes(num_imgs), binwidth = 10, fill
="purple",col ="black")

grid.arrange(p1, p2, p3, p4, p5, p6, nrow =2, ncol =3)
```




### 3.3 Bivariate Analysis - Checking effect of days and news type on shares

*#Taking log of shares column to rescale column for visualizations*

```
articles$logshares=log(articles$shares)
```

```
#Checking if publishing days make a difference
articles$news_day[articles$weekday_is_monday==1] <- "Monday"
articles$news_day[articles$weekday_is_tuesday==1] <- "Tuesday"
articles$news_day[articles$weekday_is_wednesday==1] <- "Wednesday"
articles$news_day[articles$weekday_is_thursday==1] <- "Thursday"
articles$news_day[articles$weekday_is_friday==1] <- "Friday"
articles$news_day[articles$weekday_is_saturday==1] <- "Saturday"
articles$news_day[articles$weekday_is_sunday==1] <- "Sunday"
#Check
p1 <- ggplot(articles, aes(as.factor(news_day), logshares))
p1 + geom_boxplot()
```




```
#Checking if publishing topics make a difference
articles$news_type[articles$data_channel_is_lifestyle==1] <- "Lifestyle"
articles$news_type[articles$data_channel_is_entertainment==1] <-
"Lifestyle"
articles$news_type[articles$ data_channel_is_bus==1] <- "Business"
articles$news_type[articles$data_channel_is_socmed==1] <- "Social_Media"
```

```

articles$news_type[articles$data_channel_is_tech==1] <- "Technology"
articles$news_type[articles$data_channel_is_world==1] <- "World"

p2 <- ggplot(articles, aes(as.factor(news_type), logshares))
p2 + geom_boxplot()

```



## 4. Data Preparation

### 4.1 Checking for missing values and duplicates

*#missing values check*

```
sapply(articles1,function(x) sum(is.na(x)))
```

```
##          url                                timedelta
##          0                                0
##      n_tokens_title          n_tokens_content
##          0                                0
##      n_unique_tokens          n_non_stop_words
##          0                                0
##  n_non_stop_unique_tokens          num_hrefs
##          0                                0
##      num_self_hrefs          num_imgs
##          0                                0
##      num_videos          average_token_length
##          0                                0
##      num_keywords  data_channel_is_lifestyle
##          0                                0
## data_channel_is_entertainment  data_channel_is_bus
##          0                                0
##      data_channel_is_socmed  data_channel_is_tech
##          0                                0
##      data_channel_is_world          kw_min_min
##          0                                0
##          kw_max_min          kw_avg_min
##          0                                0
##          kw_min_max          kw_max_max
##          0                                0
##          kw_avg_max          kw_min_avg
##          0                                0
##          kw_max_avg          kw_avg_avg
##          0                                0
##      self_reference_min_shares  self_reference_max_shares
##          0                                0
##  self_reference_avg_sharess          weekday_is_monday
##          0                                0
##      weekday_is_tuesday          weekday_is_wednesday
##          0                                0
##      weekday_is_thursday          weekday_is_friday
##          0                                0
##      weekday_is_saturday          weekday_is_sunday
##          0                                0
##          is_weekend          LDA_00
##          0                                0
##          LDA_01          LDA_02
##          0                                0
##          LDA_03          LDA_04
##          0                                0
```



```
##          global_subjectivity    global_sentiment_polarity
##                      0                      0
##  global_rate_positive_words    global_rate_negative_words
##                      0                      0
##          rate_positive_words    rate_negative_words
##                      0                      0
##          avg_positive_polarity    min_positive_polarity
##                      0                      0
##          max_positive_polarity    avg_negative_polarity
##                      0                      0
##          min_negative_polarity    max_negative_polarity
##                      0                      0
##          title_subjectivity    title_sentiment_polarity
##                      0                      0
##          abs_title_subjectivity    abs_title_sentiment_polarity
##                      0                      0
##                      popular
##                      0
```

*#duplicates check*

```
articles1[duplicated(articles1)]
```

```
## data frame with 0 columns and 39644 rows
```

#### 4.2 Dropping redundant features

*#removing all day leaving only is\_weekend to avoid repetition*

```
articles2 <- subset( articles1, select = -c(weekday_is_monday,
weekday_is_tuesday, weekday_is_wednesday,
                                         weekday_is_thursday,
weekday_is_friday, weekday_is_saturday,
                                         weekday_is_sunday) )
```

*#removing other non\_informative features*


```
articles3 <- subset( articles2, select = -c(url, timedelta ) )
```

#### 4.3 Checking for and removing highly correlated features

```
Cor <- round(cor(articles3),2)
```

```
corrplot(Cor, type="lower",method ="color", title = "Correlation Plot",
          mar=c(0,1,1,1), tl.cex= 0.65, outline= T, tl.col= rgb(0, 0, 0))
```

Correlation Plot



```
#Setting correlation cutoff
highlyCorrelated <- findCorrelation(Cor, cutoff = 0.7)
highlyCorCol <- colnames(articles2)[highlyCorRelated]
highlyCorCol

## [1] "global_rate_positive_words" "kw_min_avg"
## [3] "data_channel_is_socmed" "max_positive_polarity"
## [5] "global_rate_negative_words" "LDA_02"
## [7] "self_reference_avg_shares" "title_sentiment_polarity"
## [9] "data_channel_is_tech" "kw_avg_avg"
## [11] "data_channel_is_world" "kw_max_avg"
## [13] "n_tokens_title" "n_tokens_content"

#removing multicollinear variables
articles3 <- articles2[, -which(colnames(articles2) %in% highlyCorCol)]
dim(articles3)

## [1] 39644 38
```

## 5 Modelling

3 Models will be used to evaluate this classification task. First the dataset will be split into Training and Test set

```

#To achieve reproducible model; set the random seed number
set.seed(100)

# Data is split into training and test set in a 80:20 ratio
TrainingIndex <- createDataPartition(articles3$popular, p=0.75, list = FALSE)

TrainingSet <- articles3[TrainingIndex,]# Training Set
TestingSet <- articles3[-TrainingIndex,]# Test Set

```

## 5.1 Logistic Regression

```

modell1 <- glm(popular~.,family=binomial(link='logit'),data = TrainingSet,
maxit = 1000 )
summary(modell1)

```

```

##
## Call:
## glm(formula = popular ~ ., family = binomial(link = "logit"),
##      data = TrainingSet, maxit = 1000)
##
## Coefficients:
##

```

	Estimate	Std. Error	z value	Pr(> z )	
## (Intercept)	9.441e-02	1.416e-01	0.667	0.504846	
## n_tokens_title	-3.682e-03	6.043e-03	-0.609	0.542319	
## n_tokens_content	1.615e-04	3.800e-05	4.251	2.13e-05	***
## n_non_stop_unique_tokens	5.274e-03	8.359e-03	0.631	0.528135	
## num_hrefs	1.293e-02	1.510e-03	8.560	< 2e-16	***
## num_self_hrefs	-2.596e-02	3.797e-03	-6.837	8.07e-12	***
## num_imgs	6.194e-03	1.780e-03	3.479	0.000503	***
## num_videos	3.558e-03	3.357e-03	1.060	0.289194	
## average_token_length	-1.398e-01	2.183e-02	-6.407	1.48e-10	***
## num_keywords	6.578e-02	7.818e-03	8.414	< 2e-16	***
## data_channel_is_lifestyle	-1.263e-02	7.105e-02	-0.178	0.858949	
## data_channel_is_entertainment	-5.017e-01	4.655e-02	-10.777	< 2e-16	***
## data_channel_is_bus	-1.784e-01	5.897e-02	-3.025	0.002482	**
## data_channel_is_socmed	1.013e+00	6.869e-02	14.741	< 2e-16	***
## data_channel_is_tech	2.904e-01	5.612e-02	5.174	2.29e-07	***
## kw_avg_min	-3.526e-05	2.969e-05	-1.187	0.235070	
## kw_min_max	-1.149e-06	2.451e-07	-4.689	2.75e-06	***
## kw_max_max	-6.363e-07	8.125e-08	-7.831	4.84e-15	***
## kw_avg_max	5.268e-07	1.646e-07	3.200	0.001377	**
## kw_min_avg	1.450e-04	1.272e-05	11.398	< 2e-16	***
## kw_max_avg	1.919e-05	3.615e-06	5.310	1.10e-07	***
## self_reference_avg_sharess	7.165e-06	9.636e-07	7.436	1.04e-13	***
## is_weekend	8.783e-01	3.933e-02	22.332	< 2e-16	***
## LDA_01	-4.471e-01	9.266e-02	-4.825	1.40e-06	***
## LDA_02	-1.206e+00	8.228e-02	-14.660	< 2e-16	***
## LDA_03	-2.668e-01	8.340e-02	-3.200	0.001377	**
## global_subjectivity	1.220e+00	1.744e-01	6.996	2.64e-12	***
## global_sentiment_polarity	-7.266e-02	3.429e-01	-0.212	0.832190	
## global_rate_positive_words	-1.477e+00	1.263e+00	-1.169	0.242345	

```

## global_rate_negative_words      2.327e-02  2.026e+00   0.011 0.990834
## avg_positive_polarity            2.177e-02  2.850e-01   0.076 0.939109
## min_positive_polarity            -9.397e-01  2.388e-01  -3.935 8.32e-05 ***
## max_positive_polarity            -5.861e-02  8.966e-02  -0.654 0.513273
## avg_negative_polarity            2.611e-01  1.786e-01   1.462 0.143692
## max_negative_polarity            -1.652e-01  1.917e-01  -0.861 0.388985
## title_subjectivity               1.163e-01  4.526e-02   2.569 0.010212 *
## title_sentiment_polarity         2.063e-01  5.016e-02   4.113 3.91e-05 ***
## abs_title_subjectivity           1.951e-01  7.761e-02   2.514 0.011950 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 41090  on 29732  degrees of freedom
## Residual deviance: 37792  on 29695  degrees of freedom
## AIC: 37868
##
## Number of Fisher Scoring iterations: 5

#calculating errors
#test error
cut <- 0.5

yhat = (predict(model1,TrainingSet,type="response")>cut)
tr.err = mean(TrainingSet$popular != yhat)
tr.err

## [1] 0.3586251

# calculate the testing error in a similar manner to the training error
yhat = (predict(model1,TestingSet,type="response")>cut)
te.err = mean(TestingSet$popular != yhat)
print(te.err)

## [1] 0.3507214

#print(predict(cls,test,type="response")>cut)

# calculation of Naive predictor error rate where cut = 1

# so the Naive predictor will simply predict all customers stay...
# so it will make errors on each customer that leaves the bank

trN.err <- mean(!TrainingSet$popular)
teN.err <- mean(!TestingSet$popular)


print(paste("Naive train error",trN.err))

## [1] "Naive train error 0.467157703561699"

```

```
print(paste("Naive test error",teN.err))  
## [1] "Naive test error 0.4641307637978"  
hist(predict(model1,TestingSet,type="response"))
```

**Histogram of predict(model1, TestingSet, type = "response")**



```

# Prediction on TestingSet using Logistic Regression
prediction <- predict(model1, TestingSet, type = "response")
head(prediction)

##           3           8           9          12          17          18
## 0.5665468 0.6467499 0.5943744 0.5725975 0.5169225 0.5029694

#Assigning probabilities - If prediction exceeds threshold of 0.5, 1 else 0
prediction <- ifelse(prediction > 0.5, 1, 0)
head(prediction)

##  3  8  9 12 17 18
##  1  1  1  1  1  1

#Computing confusion matrix values
confusionMatrix(factor(TestingSet$popular), factor(prediction), mode
='everything', positive = "0")

## Confusion Matrix and Statistics
##
##              Reference
## Prediction    0    1
##              0 2729 1871
##              1 1605 3706
##
##              Accuracy : 0.6493
##              95% CI : (0.6398, 0.6587)
##              No Information Rate : 0.5627
##              P-Value [Acc > NIR] : < 2.2e-16
##
##              Kappa : 0.2922
##
##  Mcnemar's Test P-Value : 6.965e-06
##
##              Sensitivity : 0.6297
##              Specificity : 0.6645
##              Pos Pred Value : 0.5933
##              Neg Pred Value : 0.6978
##              Precision : 0.5933
##              Recall : 0.6297
##              F1 : 0.6109
##              Prevalence : 0.4373
##              Detection Rate : 0.2754
##              Detection Prevalence : 0.4641
##              Balanced Accuracy : 0.6471
##
##              'Positive' Class : 0
##

```

## 5.2 Decision Trees

3 cp partitions 0.01, 0.001 and 0.00001 will be used for prediction.

### 5.2.1 Tree 1 $cp = 0.01$

```
tree1 <- rpart(popular~., method = 'class', data = TrainingSet, control =  
rpart.control(cp = 0.01))
```

*#using tree1 for predicting test set*

```
test_prediction1 <- predict(tree1, TestingSet, type = 'class')  
head(test_prediction1)
```

```
## 3 8 9 12 17 18  
## 0 1 1 0 0 0  
## Levels: 0 1
```

*# converting TestingSet\$popular to factor*

```
popular_factor <- as.factor(TestingSet$popular)
```

*#confusion matrix for tree1*

```
cfm1 <- confusionMatrix(test_prediction1, popular_factor)  
cfm1
```

```
## Confusion Matrix and Statistics
```

```
##
```

```
##           Reference
```

```
## Prediction    0    1
```

```
##           0 2285 1289
```

```
##           1 2315 4022
```

```
##
```

```
##           Accuracy : 0.6364
```

```
##           95% CI : (0.6268, 0.6458)
```

```
## No Information Rate : 0.5359
```

```
## P-Value [Acc > NIR] : < 2.2e-16
```

```
##
```

```
##           Kappa : 0.2579
```

```
##
```

```
## McNemar's Test P-Value : < 2.2e-16
```

```
##
```

```
##           Sensitivity : 0.4967
```

```
##           Specificity : 0.7573
```

```
## Pos Pred Value : 0.6393
```

```
## Neg Pred Value : 0.6347
```

```
## Prevalence : 0.4641
```

```
## Detection Rate : 0.2306
```

```
## Detection Prevalence : 0.3606
```

```
## Balanced Accuracy : 0.6270
```

```
##
```

```
## 'Positive' Class : 0
```

```
##
```

```

#error rate for tree1
error_rate1 <- 1 - cfm1$overall["Accuracy"]
error_rate1

## Accuracy
## 0.3636364

```

### 5.2.2 Tree 2 $cp = 0.001$

```

tree2 <- rpart(popular~., method = 'class', data = TrainingSet, control =
rpart.control(cp = 0.001))

```

```

#using tree2 for predicting test set
test_prediction2 <- predict(tree2, TestingSet, type = 'class')
head(test_prediction2)

```

```

## 3 8 9 12 17 18
## 0 1 1 0 1 0
## Levels: 0 1

```

```

#confusion matrix for tree1
cfm2 <- confusionMatrix(test_prediction2, popular_factor)
cfm2

```

```

## Confusion Matrix and Statistics
##
##              Reference
## Prediction    0    1
##              0 2722 1579
##              1 1878 3732
##
##              Accuracy : 0.6512
##              95% CI : (0.6417, 0.6606)
##              No Information Rate : 0.5359
##              P-Value [Acc > NIR] : < 2.2e-16
##
##              Kappa : 0.2957
##
##              Mcnemar's Test P-Value : 4.013e-07
##
##              Sensitivity : 0.5917
##              Specificity : 0.7027
##              Pos Pred Value : 0.6329
##              Neg Pred Value : 0.6652
##              Prevalence : 0.4641
##              Detection Rate : 0.2746
##              Detection Prevalence : 0.4340
##              Balanced Accuracy : 0.6472
##
##              'Positive' Class : 0
##

```



```

#error rate for tree1
error_rate1 <- 1 - cfm2$overall["Accuracy"]
error_rate1

## Accuracy
## 0.3488044

```

### 5.2.3 Tree 3, cp = 0.00001

```

tree3 <- rpart(popular~., method = 'class', data = TrainingSet, control =
rpart.control(cp = 0.00001))

```

#### #using tree1 for predicting test set

```

test_prediction3 <- predict(tree3, TestingSet, type = 'class')
head(test_prediction3)

```

```

## 3 8 9 12 17 18
## 0 1 0 1 0 0
## Levels: 0 1

```

#### #confusion matrix for tree1

```

cfm3 <- confusionMatrix(test_prediction3, popular_factor)
cfm3

```

#### ## Confusion Matrix and Statistics

```

##
##           Reference
## Prediction    0    1
##           0 2559 2062
##           1 2041 3249
##
##           Accuracy : 0.586
##           95% CI : (0.5762, 0.5957)
##           No Information Rate : 0.5359
##           P-Value [Acc > NIR] : <2e-16
##
##           Kappa : 0.168
##
##           McNemar's Test P-Value : 0.7549
##
##           Sensitivity : 0.5563
##           Specificity : 0.6117
##           Pos Pred Value : 0.5538
##           Neg Pred Value : 0.6142
##           Prevalence : 0.4641
##           Detection Rate : 0.2582
##           Detection Prevalence : 0.4662
##           Balanced Accuracy : 0.5840
##
##           'Positive' Class : 0
##

```

```

#error rate for tree1
error_rate3 <- 1 - cfm3$overall["Accuracy"]
error_rate3

## Accuracy
## 0.4139845

```

### 5.3 Model 3 - Random Forest

*#First step in running rf is converting target variable to factor*  
 TrainingSet\$popular <- as.factor(TrainingSet\$popular)

#### 5.3.1 - Random Forest ntree = 100

```

# Assuming your data frame is called 'df' and the target variable is 'target'
rf_model <- randomForest(popular~ ., data = TrainingSet, ntree = 100)
rf_model

##
## Call:
## randomForest(formula = popular ~ ., data = TrainingSet, ntree = 100)
##              Type of random forest: classification
##              Number of trees: 100
## No. of variables tried at each split: 6
##
## OOB estimate of error rate: 34.47%
## Confusion matrix:
##      0      1 class.error
## 0 8246  5644   0.4063355
## 1 4606 11237   0.2907278

rf_predictions <- predict(rf_model, TestingSet)
head(rf_predictions)

##  3  8  9 12 17 18
##  0  0  1  0  0  0
## Levels: 0 1

#confusion matrix for rf_model
cf_rf <- confusionMatrix(rf_predictions, popular_factor)
cf_rf

## Confusion Matrix and Statistics
##
##              Reference
## Prediction      0      1
##              0 2731 1486
##              1 1869 3825
##
##              Accuracy : 0.6615
##              95% CI : (0.6521, 0.6708)
##              No Information Rate : 0.5359
##              P-Value [Acc > NIR] : < 2.2e-16

```

```
##
##          Kappa : 0.3157
##
## Mcnemar's Test P-Value : 4.252e-11
##
##          Sensitivity : 0.5937
##          Specificity : 0.7202
##          Pos Pred Value : 0.6476
##          Neg Pred Value : 0.6718
##          Prevalence : 0.4641
##          Detection Rate : 0.2756
##          Detection Prevalence : 0.4255
##          Balanced Accuracy : 0.6569
##
##          'Positive' Class : 0
##

#error rate for tree1
rf_error_rate <- 1 - cf_rf$overall["Accuracy"]
rf_error_rate

## Accuracy
## 0.3385128
```

### 5.3.2 Random Forest ntree = 500

```
rf_model2 <- randomForest(popular~ ., data = TrainingSet, ntree = 500,
importance = TRUE)
rf_model2

##
## Call:
## randomForest(formula = popular ~ ., data = TrainingSet, ntree = 500,
importance = TRUE)
##          Type of random forest: classification
##          Number of trees: 500
## No. of variables tried at each split: 6
##
##          OOB estimate of  error rate: 33.54%
## Confusion matrix:
##      0      1 class.error
## 0 8254  5636   0.4057595
## 1 4336 11507   0.2736855

rf2_predictions <- predict(rf_model2, TestingSet)
head(rf2_predictions)

##  3  8  9 12 17 18
##  0  0  0  0  0  0
## Levels: 0 1
```

```

#confusion matrix for rf_model2
cf_rf2 <- confusionMatrix(rf2_predictions, popular_factor)
cf_rf2

## Confusion Matrix and Statistics
##
##              Reference
## Prediction    0    1
##              0 2761 1455
##              1 1839 3856
##
##              Accuracy : 0.6676
##              95% CI : (0.6583, 0.6769)
##              No Information Rate : 0.5359
##              P-Value [Acc > NIR] : < 2.2e-16
##
##              Kappa : 0.3281
##
##  Mcnemar's Test P-Value : 2.502e-11
##
##              Sensitivity : 0.6002
##              Specificity : 0.7260
##              Pos Pred Value : 0.6549
##              Neg Pred Value : 0.6771
##              Prevalence : 0.4641
##              Detection Rate : 0.2786
##              Detection Prevalence : 0.4254
##              Balanced Accuracy : 0.6631
##
##              'Positive' Class : 0
##

#error rate for rf2
rf2_error_rate <- 1 - cf_rf2$overall["Accuracy"]
rf2_error_rate

## Accuracy
## 0.332358

```

#### 5.4 Calculating Feature Importance using Random Forest

```

importance_values <- importance(rf_model2)
importance_values

```


	0	1	MeanDecreaseAccuracy
n_tokens_title	4.363174	3.33854834	5.529538
n_tokens_content	20.794649	18.34319220	34.415810
n_non_stop_unique_tokens	20.778322	13.74512270	30.688867
num_hrefs	11.719759	15.51879073	22.494046
num_self_hrefs	10.780497	4.40122692	12.167521
num_imgs	26.219943	2.01084468	23.400708
num_videos	17.506425	-1.19419626	12.954985

## average_token_length	28.241977	-0.95029873	24.282404
## num_keywords	11.891724	5.83443767	15.098458
## data_channel_is_lifestyle	13.092010	-10.31218676	4.701038
## data_channel_is_entertainment	33.641036	16.09649599	36.376499
## data_channel_is_bus	15.998257	-1.67418147	15.362631
## data_channel_is_socmed	17.842549	35.25916542	36.405930
## data_channel_is_tech	21.259648	16.63668362	28.211243
## kw_avg_min	17.212975	13.35841256	23.555205
## kw_min_max	3.061054	22.65037524	28.356639
## kw_max_max	11.697241	16.47082765	22.876947
## kw_avg_max	23.317106	2.21467488	23.915655
## kw_min_avg	27.633594	19.41837586	43.094963
## kw_max_avg	37.565745	34.61000478	52.372030
## self_reference_avg_sharess	34.550917	33.74759922	51.720217
## is_weekend	52.116303	24.22507614	51.283820
## LDA_01	18.900924	15.88545286	28.738013
## LDA_02	36.369859	9.95840729	42.441231
## LDA_03	21.964375	5.08035386	23.209032
## global_subjectivity	12.352345	14.63540257	21.231503
## global_sentiment_polarity	19.630304	9.49759093	24.411006
## global_rate_positive_words	17.801403	11.77461801	24.273712
## global_rate_negative_words	12.073887	7.14262865	16.323985
## avg_positive_polarity	4.585040	13.41228729	14.774362
## min_positive_polarity	5.047227	17.74365088	20.518451
## max_positive_polarity	6.508848	10.54839435	14.078436
## avg_negative_polarity	9.085544	7.46072631	13.266155
## max_negative_polarity	8.573425	5.53425953	11.396534
## title_subjectivity	10.017977	3.68968268	10.705089
## title_sentiment_polarity	11.321232	0.95380545	8.744615
## abs_title_subjectivity	8.840689	-0.07520549	6.539739
##	MeanDecreaseGini		
## n_tokens_title	332.22283		
## n_tokens_content	564.58310		
## n_non_stop_unique_tokens	583.52040		
## num_hrefs	422.74092		
## num_self_hrefs	279.25186		
## num_imgs	292.07579		
## num_videos	166.74953		
## average_token_length	570.51392		
## num_keywords	217.50046		
## data_channel_is_lifestyle	33.67307		
## data_channel_is_entertainment	165.38118		
## data_channel_is_bus	54.26601		
## data_channel_is_socmed	102.69582		
## data_channel_is_tech	118.90734		
## kw_avg_min	634.55194		
## kw_min_max	347.37658		
## kw_max_max	147.74721		
## kw_avg_max	595.58835		
## kw_min_avg	496.82662		

```
## kw_max_avg          903.40462
## self_reference_avg_shares 752.87178
## is_weekend          228.62720
## LDA_01              617.16050
## LDA_02              755.81093
## LDA_03              577.19293
## global_subjectivity  574.89423
## global_sentiment_polarity 546.93224
## global_rate_positive_words 554.40838
## global_rate_negative_words 514.56847
## avg_positive_polarity 526.78724
## min_positive_polarity 256.15986
## max_positive_polarity 211.59787
## avg_negative_polarity 507.31990
## max_negative_polarity 287.07209
## title_subjectivity  276.33602
## title_sentiment_polarity 325.06909
## abs_title_subjectivity 244.63308
```

```
varImpPlot(rf_model2)
```

rf\_model2



## 6 Model Evaluation

Models are evaluated using accuracy from confusion matrix, testing error and also AUC score.

### 6.1 Table of Results

*# Create a new table with some sample data*

```
Model_Comparison <- data.frame(  
  Model = c("Logistic Regression", "Decison Trees", "Random Forest"),  
  Accuracy = c(0.644, 0.645, 0.668),  
  TestingError = c(0.356, 0.355, 0.332),  
  Sensitivity = c(0.625, 0.585, 0.595),  
  Specificity = c(0.659, 0.702, 0.732))
```

*# Display the new table*

```
print(Model_Comparison)
```

```
##           Model Accuracy TestingError Sensitivity Specificity  
## 1 Logistic Regression    0.644         0.356      0.625      0.659  
## 2      Decison Trees    0.645         0.355      0.585      0.702  
## 3      Random Forest    0.668         0.332      0.595      0.732
```

### ###6.2 Plotting ROC Curves

*#converting prediction scores data type before plotting curves*

```
test_prediction2 <- as.numeric(test_prediction2)
```

```
rf2_predictions <- as.numeric(rf2_predictions)
```

*#creating the ROC function*

```
glm_roc_curve <- roc(TestingSet$popular, prediction)
```

```
## Setting levels: control = 0, case = 1
```

```
## Setting direction: controls < cases
```

```
tree_roc_curve <- roc(TestingSet$popular, test_prediction2)
```

```
## Setting levels: control = 0, case = 1
```

```
## Setting direction: controls < cases
```

```
rf_roc_curve <- roc(TestingSet$popular, rf2_predictions)
```

```
## Setting levels: control = 0, case = 1
```

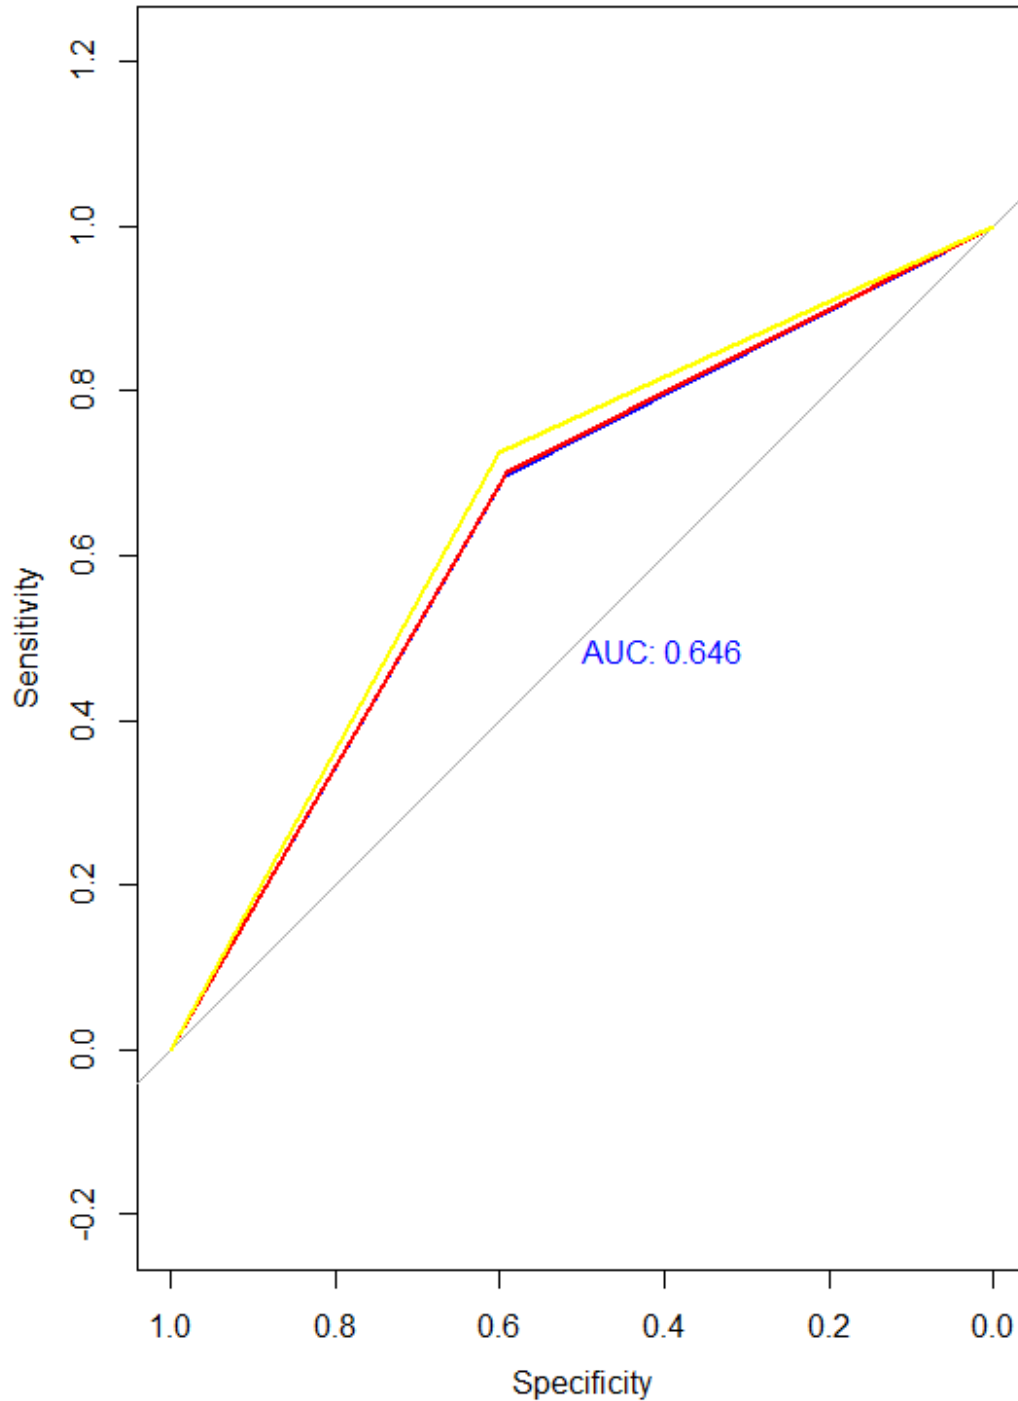
```
## Setting direction: controls < cases
```

*# Plotting ROC Curves*

```
plot(glm_roc_curve, col = "blue", print.auc = TRUE)
```

```
plot(tree_roc_curve, col = "red", add = TRUE)
```

```
plot(rf_roc_curve, col = "yellow", add = TRUE)
```



```
#calculating AUC curves of models# Calculate ROC and AUC using pROC
glm_score <- print(paste('glm roc_roc_curve score is',auc(glm_roc_curve)))
## [1] "glm roc_roc_curve score is 0.645528947303791"
tree_score <- print(paste('tree_roc_curve score is',auc(tree_roc_curve)))
```



```
## [1] "tree_roc_curve score is 0.647215827691502"  
rf2_score <- print(paste('rf_roc_curve score is', auc(rf_roc_curve)))  
## [1] "rf_roc_curve score is 0.663128842517171"
```

Proceeding with Random Forest because it has the highest accuracy 66.8% and AUC score of 0.663